

Technical Documentation: Daniel Rada AI Assistant Table of Contents Project Overview

System Architecture

Core Components

Support Scripts

Workflow

Key Mechanisms

Technical Considerations

Future Improvements

Troubleshooting

1. Project Overview Personal AI assistant system using RAG (Retrieval-Augmented Generation) to answer questions about Daniel Rada's professional experience. Combines document processing, language models, and vector storage.

Key Tools:

Backend: FastAPI (Python)

Frontend: Flask (Python)

LLM Model: Mistral-7B (quantized GGUF)

Vector Store: ChromaDB

Embeddings: Multilingual Sentence Transformers

Memory: Redis

PDF Processing: PyPDF2 + Tesseract OCR

2. System Architecture Diagram Code

3. Core Components 3.1 Document Processing (pdf_to_text.py) Function: Convert CVs/PDFs to structured text

Features:

Automatic detection of scanned PDFs (OCR with Tesseract)

Parallel processing with ThreadPool

Text cleaning with regex

Usage:

```
bash python pdf_to_text.py --input pdfs --output data
```

3.2 Embedding Creation (ingest.py) Function: Generate embeddings and store in ChromaDB

Workflow:

Load TXT documents

Split text into chunks (800 tokens with 150 overlap)

Generate embeddings with multilingual model

Store in persistent ChromaDB

Model: sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2

3.3 Backend API (api.py) Endpoints:

POST /ask: Process questions using RAG

POST /feedback: Receive response corrections

GET /health: Check system status

Features:

Conversation memory with Redis

Session management

Detailed logging

Dependencies:

```
python vector_store = Chroma(persist_directory="embeddings")
```

3.4 RAG Model (model.py) Components:

Prompt Template: Strict bilingual structure (ES/EN)

LLM Model: Mistral-7B (4-bit quantized)

Configuration:

python LlamaCpp(model_path=".../mistral-7b.Q4_K_M.gguf", n_gpu_layers=33, temperature=0.3) Retrieval: MMR (Maximal Marginal Relevance) search with 4 relevant documents

3.5 Web Interface (web_app.py) Technology: Flask + Jinja2

Features:

Unique session ID generation

Backend API communication

Integrated feedback system

Structure:

text /web |— templates/ | |— index.html |— static/ |— web_app.py

4. Support Scripts Script Function rename.sh Remove spaces from filenames remove_dashes.sh Remove hyphens from filenames update_ai.sh Update embeddings and restart services compatibility_check.py Verify dependency versions

5. Workflow Installation: Install dependencies:

bash pip install -r requirements.txt Download Mistral-7B model:

bash ./scripts/download_model.sh Start Redis:

bash docker run -p 6379:6379 redis Initial Processing: Diagram Code Execution: bash

Backend (API)

python api.py

Frontend (Web)

python web_app.py

6. Key Mechanisms Memory Management Stores conversation history per session in Redis

Key format: danielai:<session_id>

Retention: 30 days (configurable)

Feedback System Stores corrections in CSV:

text question,response,correction Path: feedback_data.csv

Language Handling Strict rules in prompt:

text IF spanish_question → spanish_response IF english_question → english_response

7. Technical Considerations Hardware Requirements:

Recommended GPU (≥ 8 GB VRAM)

Minimum 16GB RAM

Storage: 10GB+ (models + embeddings)

Optimizations:

4-bit model quantization

Batch document processing

Redis shared memory usage

Security:

Sessions don't store personal data

Redis not publicly exposed

Input validation with Pydantic

8. Future Improvements Implement fine-tuning with feedback data

Add JWT authentication

Priority system for key documents

Conversation monitoring dashboard

9. Troubleshooting Issue: Model fails to load Solution: Verify model path in model.py

Issue: ChromaDB not persisting Solution: Check permissions for embeddings/ directory

Issue: Responses in wrong language Solution: Validate prompt template in model.py

bash

Verify compatibility:

python compatibility_check.py Note: This documentation covers system version 1.0. For updates, refer to the official repository.